

# Kubernetes Workshop

The Basics





# **Ansgar Sachs**

## **Senior Software Developer**

### **CGM Software GmbH**

ansgar.sa@gmail.com

# What do you already know and what's your goal of this workshop?

## JAY ESDOT

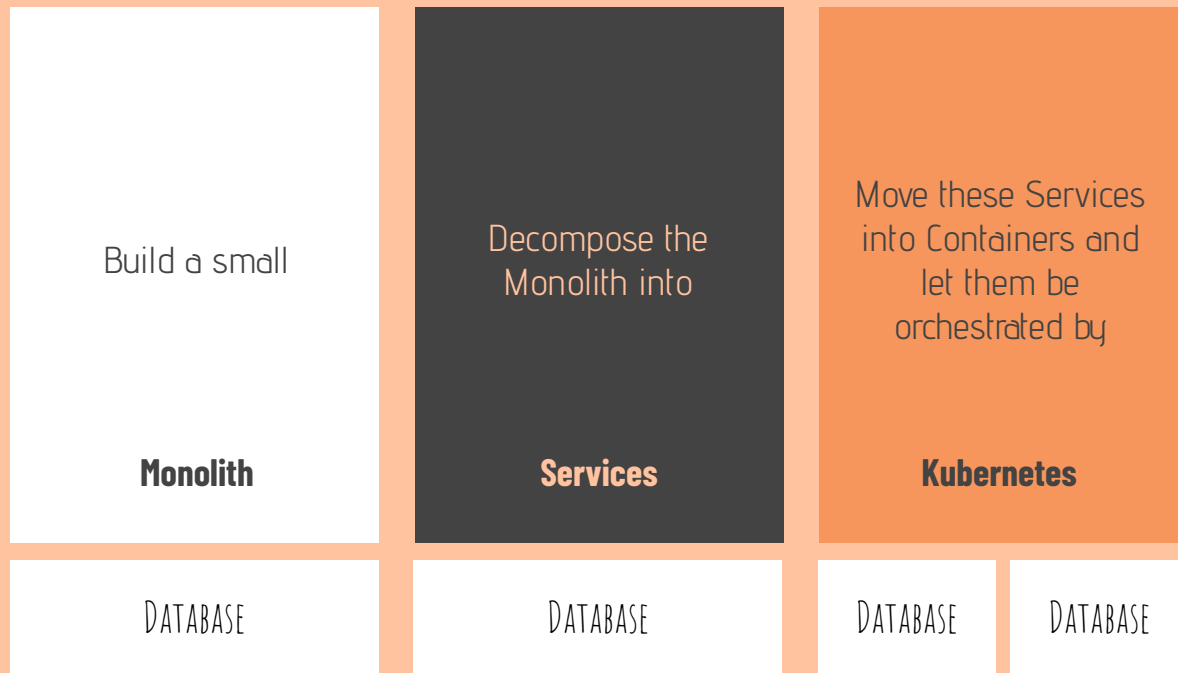
Knows about Docker but has  
never used Kubernetes

## JAYLINA TEDOT

Has seen Docker a while  
ago but isn't really  
comfortable with it



# Goal of this workshop



**YOU GET A CLUSTER AND YOU GET A CLUSTER**



**EVERYBODY GETS A CLUSTER**

# Build Pairs

If you are already experienced, then search for someone with less knowledge about containers and vice versa

Check out:

[HTTPS://GITHUB.COM/ANSGARS/KUBERNETES-WORKSHOP](https://github.com/ansgars/kubernetes-workshop)

# But before...

Here are some Spring fundamentals you should know

spring



```
@SpringBootApplication
public class MyAwesomeApp {

    public static void main(String[] args) {
        SpringApplication.run(MyAwesomeApp.class, args);
    }

    @Bean(name = "someService")
    public SomeService someService() {
        return new SomeService();
    }
}
```

# Spring Basics 1



# Spring Basics 2

```
@RestController
@RequestMapping(value = "/iamroute")
public class SomeController {
    Logger log = LoggerFactory.getLogger(SomeController.class);

    @Resource(name = "someService")
    private SomeService someService;

    @PostMapping(produces = {"application/json"})
    public ResponseEntity<Sth> addSth(@RequestBody Blub blub) {
        // ...
        return new ResponseEntity<>(blob, HttpStatus.CREATED);
    }
}
```

# Task 1

Build a monolithic spring app, that

- accepts new patients (id, firstName, lastName, age)
- accepts new prescriptions (id, name, dose, price)
- returns the cumulated costs of a patient's prescriptions
- use an in memory data store

Use a client such as Postman to validate your application **and** run the integration tests

```
cd integration-tests && mvn clean  
verify
```



# Served Routes

```
POST /patients {  
  "firstName": "...",  
  "lastName": "...",  
  "age": 123  
}
```

Returns 201 Created { "id": "..." }

```
POST /prescriptions {  
  "patientId": "...",  
  "name": "...",  
  "price": 12.12  
}
```

Returns 201 Created { "id": "..." }

Returns 422 if patient doesn't exist

```
GET /patients/<id>/costs
```

```
Returns 200 OK {  
  "patientId": "...",  
  "costs": 45.27  
}
```

Returns 404 Not found if patient does not exist

```
GET /patients
```

```
Returns 200 OK {  
  "patients": ["id1", "id2"]  
}
```

# Task 1: Result

If you could not finish your task

- checkout task1/final
- run the integration tests once



## Questions

- Did you create different models for the write and the read model?
- How did you manage your singleton in memory state?
- How did you cluster your controllers and actions?



# 01

## Baseline

Let's create something to iterate on

# 02

## Docker Fundamentals

Overview of the most important components

# 03

## Kubernetes 101

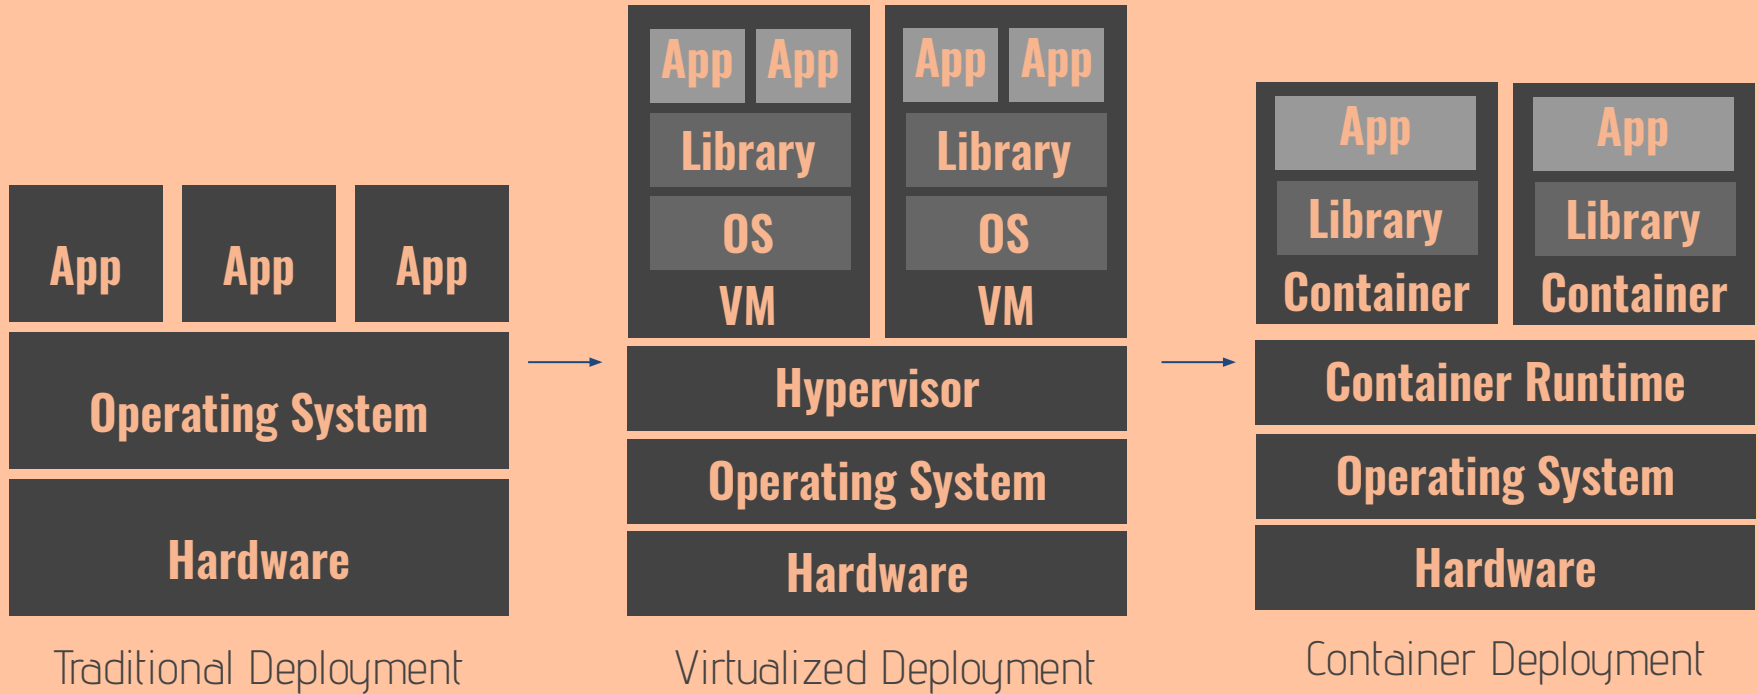
What does it take to create a mini application?

# 04

## Discussion

I understand that this topic is complex. What are the next steps?

# A brief History of Deployments



# Docker

A platform to separate your application from your infrastructure

- package and run applications in a loosely isolated environment (container)
- The container becomes the unit for distributing and testing your application
- The isolation and security allow you to run many containers simultaneously on a given host
- A registry keeps track of your application



**Container,  
Docker and  
everything else...**

# Dockerfile

```
FROM SomeBaseImage
```

```
RUN someShellCommand
```

```
ENV SOME_ENV SOME_VALUE
```

```
ADD someFile
```

```
Copy someOtherFile
```

```
EXPOSE 9999
```

```
CMD ["Some Start Command"]
```



# Container, Docker and everything else...

## Docker Client

```
docker build -t maintainer/imageName:tag .
```

```
docker pull maintainer/imageName:tag
```

```
docker run -d -p 8000:8000 --name=portainer \  
-v host:container maintainer/image
```

**Build a Docker image**

**Pull a Docker image**

**Create and run a Docker Container**

## `docker-compose.yml`

```
version: "3.7"

services:
  some-service:
    build:
      context: ./service-directory
    ports:
      - "9999:9999"
    networks:
      - a-network
    volumes:
      - a_volume:/container/path
    environment:
      - POSTGRES_USER=cgmuser
    ...
```

## `docker-compose.yml`

```
...

volumes:
  a_volume:

networks:
  some-network:
    driver: bridge
```

Container,  
Docker and  
everything else...

# Containers, Docker and everything else...

**Build a Docker Image**

**Pull a Docker Image**

**Create and run a Docker Container**

**Shutdown composed cluster**

## Docker Compose Client

**docker-compose build service-name**

**docker-compose pull service-name**

**docker-compose up -d service-name**

**docker-compose down**

## Task 2

Split up the monolithic spring app and containerize it in a composed platform

- first spring app accepts new patients and prescriptions
- second spring app returns a list of all patients or a list of all prescriptions for a given patient
- all apps are connected to a postgres database

**Note:** Merge **task2/starter** into your current branch


Use a client such as Postman to validate your application or run the integration tests



# Inside Kubernetes 03.

A brief story about the Architecture





Kubernetes does the things that the very best system administrator would do: automation, failover, centralized logging, monitoring. It takes what we've learned in the devops community and makes it the default, out of the box.

— **KELSEY HIGHTOWER**

# Why Kubernetes?

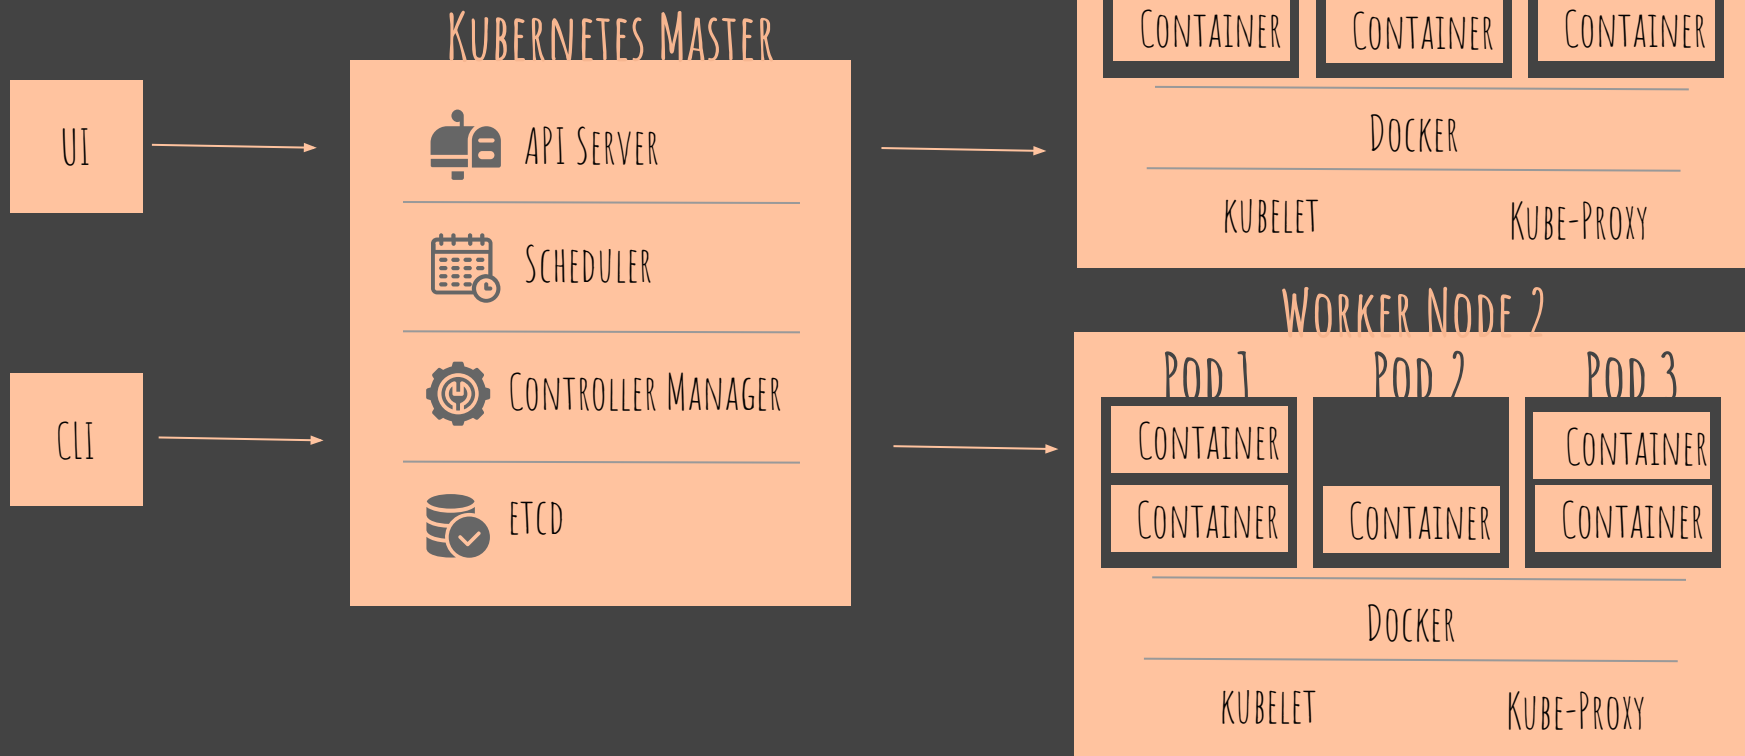
Organizations that operate on a massive scale require:

- an architecture that scales with the usage
- an architecture that scales with development teams
- an architecture that supports ongoing operations
- an architecture that supports automated development cycles



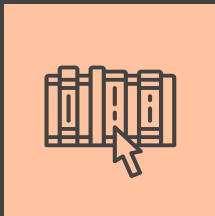
“ More than 75% of global organizations will be running containerized applications in production by 2022 “  
- see [Gartner](#)

# The k8s Architecture



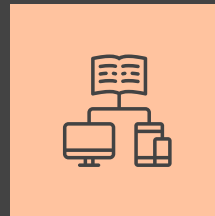


# The two big players



## Kubernetes Master

- Responsible for exposing the application program interface (API)
- Scheduled the deployments
- Manages the overall cluster



## Node

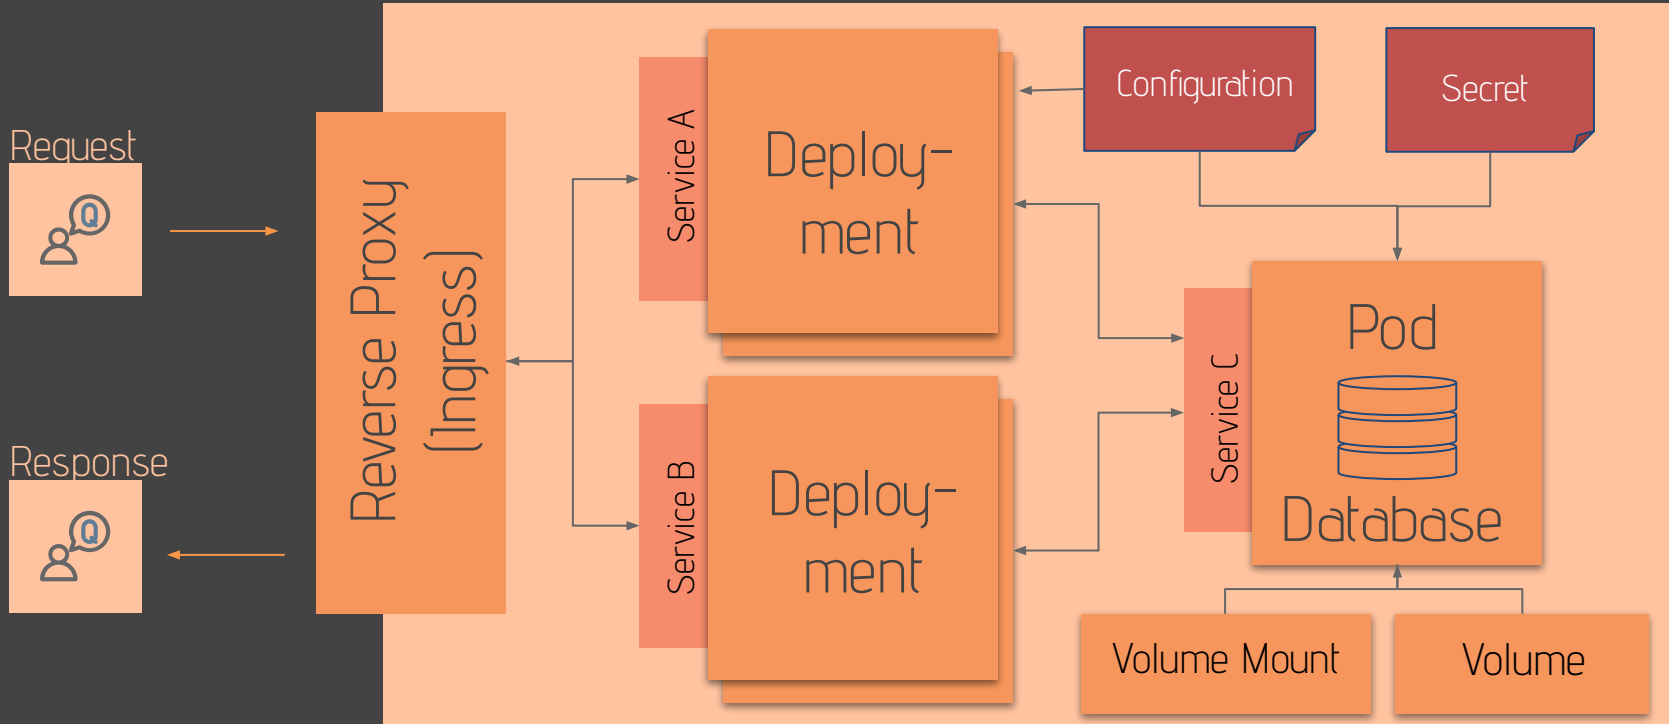
- The Workhorses of a Kubernetes Cluster
- Expose compute, networking and storage resources to applications
- Communicates via an agent with the Master

# Kubernetes Components

A RUSH THROUGH THE  
ECOSYSTEM



# Let's start with an example



```
1 apiVersion: networking.k8s.io/v1beta1
2 kind: Ingress
3 metadata:
4   name: test-ingress
5   annotations:
6     nginx.ingress.kubernetes.io/rewrite-target: /$2
7 spec:
8   rules:
9     - host: some-url.info
10      http:
11        paths:
12          - path: /testpath(/|$)(.*)
13            backend:
14              serviceName: test
15              servicePort: 9999
```

# Ingress

An API object that manages external access to the services in a cluster, typically HTTP

Further Reading:  
<https://kubernetes.io/docs/concepts/services-networking/ingress/>

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: a-service-name
5   namespace: default
6   labels:
7     app: the-service-label
8 spec:
9   ports:
10    - port: 9999
11      protocol: TCP
12   selector:
13     app: the-pod-label
```

## Service

An abstract way to expose an application running on a set of Pods as a network service.

Further Reading:  
<https://kubernetes.io/docs/concepts/services-networking/service/>

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: deployment-name
5   namespace: default
6   labels:
7     app: a-label-I-can-select
8 spec:
9   replicas: 1
10  selector:
11    matchLabels:
12      app: the-selectable-label
13  template:
14    metadata:
15      labels:
16        app: give-me-the-label-again
```

## Pod

Pods are the smallest deployable units of computing that can be created and managed in Kubernetes.

## Deployment

A Deployment provides declarative updates for Pods and ReplicaSets.

## Replica Set

A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time.

Further Reading:  
<https://kubernetes.io/docs/concepts/workloads/controllers/deployment/>

```
17     app: give-me-the-label-again
18 spec:
19   containers:
20     - name: how-i-call-the-pod
21       image: docker-image-name:<tag>
22       imagePullPolicy: IfNotPresent
23       ports:
24         - containerPort: 9999
25       env:
26         - name: SUPER_SECRET_ENV
27           valueFrom:
28             secretKeyRef:
29               name: a-secret-name
30               key: key-in-secret
```

## Pod

Pods are the smallest deployable units of computing that can be created and managed in Kubernetes.

## Deployment

A Deployment provides declarative updates for Pods and ReplicaSets.

## Replica Set

A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time.

Further Reading:

<https://kubernetes.io/docs/concepts/workloads/controllers/deployment/>

```
31     volumeMounts:
32       - mountPath: /some/path
33         name: volume-name
34       - mountPath: some-file.txt
35         name: other-volume-name
36         readOnly: true
37   volumes:
38     - name: volume-name
39       persistentVolumeClaim:
40         claimName: volume-claim-name
41     - name: other-volume-name
42       configMap:
43         name: configmap-name
```

## Pod

Pods are the smallest deployable units of computing that can be created and managed in Kubernetes.

## Deployment

A Deployment provides declarative updates for Pods and ReplicaSets.

## Replica Set

A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time.

Further Reading:  
<https://kubernetes.io/docs/concepts/workloads/controllers/deployment/>



```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: some-configmap-name
5    namespace: default
6  data:
7    somefile.sql: |
8      \connect someDatabaseName;
9
10     CREATE TABLE blub (
11         id          varchar(255) PRIMARY KEY,
12         yolo        varchar(255) NOT NULL
13     );
```

## ConfigMap

A ConfigMap is an API object used to store non-confidential data in key-value pairs. Pods can consume ConfigMaps as environment variables, command-line arguments, or as configuration files in a volume.

Further Reading:  
<https://kubernetes.io/docs/concepts/configuration/configmap/>

```
1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: some-secret-name
5 type: Opaque
6 data:
7   dbname: cGF0aWVudGRi
8   username: Y2dtdXNlcm5hbnU=
9   password: Y2dtcGFzc3dvcmQ=
```



Values are  
Base64 encoded

```
# Encode in Base64
echo -n 'username' | base64
# Decode from Base64
echo -n 'dXNlcm5hbnU=' | base64 --decode
```

## Secret

Kubernetes Secrets let you store and manage sensitive information, such as passwords, OAuth tokens, and ssh keys.

Further Reading:  
<https://kubernetes.io/docs/concepts/configuration/secret/>

```
1 kind: PersistentVolumeClaim
2 apiVersion: v1
3 metadata:
4   name: some-volume-claim-name
5   labels:
6     app: some-claim-label
7 spec:
8   accessModes:
9     - ReadWriteMany
10  resources:
11    requests:
12      storage: 1337Gi
```



There are many, many more types of volumes

## Volumes

At its core, a volume is just a directory, possibly with some data in it, which is accessible to the Containers in a Pod.

Further Reading:  
<https://kubernetes.io/docs/concepts/storage/volumes/>

# Kubernetes 101

Deploy Files  
Undeploy Files  
Get deployed Files

Inspect Files

Get logs

## Kubernetes Client

```
kubectl apply -f /path/to/directory
kubectl delete -f /path/to/directory
kubectl get deployments
services
ingress
pvc
pods
secrets
configmaps
kubectl describe deployment <name>
....
kubectl logs -f <pod-name>
```

# Task 3

Deploy all your previous built docker images in Kubernetes

- Add content to your kubernetes files in the k8s directory
- Recreate your database with a deployment, a volume, a secret, a configMap and a service
- Recreate your services with a deployment and a service file and the previous used secret file
- Use an ingress to route your traffic
- Update all application properties to the new routings

**Note:** Merge **task3/starter** into your current branch

Use a client such as Postman to validate your application and run the integration tests





# Discussion 04.

The State of Play



# SUGGESTIONS

## JAY ESDOT

Which parts of this workshop can  
be reused for my work  
environment?  
Did this workshop support me for  
future projects?

## JAYLINA TEDOT

Where can I continue learning and am  
I now able to create a production ready  
cluster?





# THANKS!

Does anyone have any questions?

[ansgar.sa@gmail.com](mailto:ansgar.sa@gmail.com)

<https://www.cgm.com/de/index.de.jsp>



# RESOURCES

The content of this presentation was stolen from

- <https://kubernetes.io/docs/concepts/>
- <https://docs.docker.com/get-started/overview/>
- <https://docs.spring.io/spring/docs/current/spring-framework-reference/index.html>
- Thousands of youtube tech talks
- my brain

# CREDITS

The following awesome resources have been used to create this presentation

- ◀ Presentation template by [Slidesgo](#)
- ◀ Icons by [Flaticon](#)
- ◀ Images & infographics by [Freepik](#)
- ◀ Author introduction slide photo created by Freepik
- ◀ Text & Image slide photo created by Freepik.com
- ◀ Big image slide photo created by Freepik.com